

effnet_rcnn_style

March 25, 2025

```
[1]: import torch
import pandas as pd
import numpy as np
from tqdm import tqdm
import wandb
```

```
[2]: wandb.
    ↪init(entity="ameyar3103-iiit-hyderabad",project="recurrent_conv_art_effnet",
    ↪config={
        "epochs": 20,
        "batch_size": 32,
        "learning_rate": 0.001,
        "model": "EFFNET_RCNN"
    })
```

wandb: Using wandb-core as the SDK backend. Please refer to <https://wandb.me/wandb-core> for more information.

wandb: Currently logged in as: ameyar3103 (ameyar3103-iiit-hyderabad) to <https://api.wandb.ai>. Use `wandb login --relogin` to force relogin

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

[2]: <wandb.sdk.wandb_run.Run at 0x714a24c95720>

0.1 Data loading

```
[3]: df_train = pd.read_csv('wikiart_csv/style_train.csv',header=None,
    ↪names=["image_path", "style_id"])
df_val = pd.read_csv('wikiart_csv/style_val.csv',header=None,
    ↪names=["image_path", "style_id"])
```

```
[4]: # get the number of classes
num_classes = 27 # from style_class.txt

[5]: # Gather input data
train_images = df_train['image_path'].values
train_labels = df_train['style_id'].values

val_images = df_val['image_path'].values
val_labels = df_val['style_id'].values

[6]: from torchvision import transforms
import cv2
```

0.2 Preprocess data and create test and train dataset

```
[7]: # create test and train dataset for dataloader

def get_image(image_path, image_size=224):
    try:
        img = cv2.imread('./wikiart/' + image_path)
        if img is None:
            raise ValueError(f"Image not loaded: ./wikiart/{image_path}")
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
        h, w, _ = img.shape
        scale = 256 / min(h, w)
        new_w = int(w * scale)
        new_h = int(h * scale)
        img_resized = cv2.resize(img, (new_w, new_h))
        start_x = (new_w - image_size) // 2
        start_y = (new_h - image_size) // 2
        img_cropped = img_resized[start_y:start_y+image_size, start_x:
↪start_x+image_size]
        img_cropped = img_cropped.astype(np.float32) / 255.0
        img_tensor = torch.from_numpy(img_cropped).permute(2, 0, 1)
        mean = torch.tensor([0.485, 0.456, 0.406]).view(3, 1, 1)
        std = torch.tensor([0.229, 0.224, 0.225]).view(3, 1, 1)
        img_tensor = (img_tensor - mean) / std
        return img_tensor
    except Exception as e:
        print(f"Error processing {image_path}: {e}")
        return torch.zeros(3, image_size, image_size)

class WikiArtDataset(torch.utils.data.Dataset):
    def __init__(self, images, labels):
        self.images = images
        self.labels = labels
```

```

def __len__(self):
    return len(self.images)

def __getitem__(self, idx):
    # image_vectors = []
    # for image in self.images:
    #     image_emb = get_image(image)
    #     image_vectors.append(image_emb)
    # image = torch.stack(image_vectors)
    image = self.images[idx]
    # label should be a one-hot encoded vector
    label = torch.zeros(num_classes)
    label[self.labels[idx]] = 1

    return image, label

train_dataset = WikiArtDataset(train_images, train_labels)
train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=16,
    ↪shuffle=True)
val_dataset = WikiArtDataset(val_images, val_labels)
val_loader = torch.utils.data.DataLoader(val_dataset, batch_size=16,
    ↪shuffle=False)

for i, (images, labels) in enumerate(train_loader):
    print(images)
    print(labels)
    break

```

```

('Post_Impressionism/paul-gauguin_and-the-gold-of-their-bodies-1901.jpg',
'Rococo/maurice-quentin-de-la-tour_portrait-of-madame-de-pompadour-1752.jpg',
'Fauvism/aldemir-martins_galo-1987.jpg', 'Baroque/adriaen-van-ostade_portrait-
of-a-boy.jpg', 'Baroque/jan-steen_merry-company-on-a-terrace-1675.jpg',
'Baroque/francisco-de-zurbaran_st-romanus-and-st-barulas-of-antioch-1638.jpg',
'Post_Impressionism/pierre-bonnard_girl-with-a-dog-in-the-park-at-grand-lemps-
also-known-as-dauphine-1900.jpg', 'Realism/gustave-courbet_the-white-
sail-1877.jpg', 'Impressionism/walter-sickert_mrs-barrett.jpg', 'Rococo/fyodor-
rokotov_portrait-of-catherine-ii-repeat-version-of-a-portrait-after-1768.jpg',
'Realism/konstantin-somov_lady-in-blue-portrait-of-the-artist-yelizaveta-
martynova-1900.jpg', 'Expressionism/martiros-saryan_in-armenia-1923.jpg',
'Realism/giovanni-boldini_portrait-of-guiseppe-verdi-1813-1901-1886.jpg',
'Impressionism/pierre-auguste-renoir_still-life-with-pomegranates.jpg',
'High_Renaissance/luca-signorelli_dante-and-virgil-entering-purgatory-1502.jpg',
'Pop_Art/claes-oldenburg_free-stamp-at-cleveland-city-hall-collaboration-with-
van-bruggen.jpg')

```

```

tensor([[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
        0., 0., 1., 0., 0., 0., 0., 0., 0.],
        [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,

```

[illegible]

```
[ ]: # CNN model
```

```
import torch.nn as nn
import torch.nn.functional as F
import torchvision.models as models

class EffNetLSTM(nn.Module):
    def __init__(self, num_classes):
        super().__init__()

        # EfficientNet-B0 backbone (outputs 1280 channels)
        effnet = models.efficientnet_b0(weights=models.EfficientNet_B0_Weights.
        IMAGENET1K_V1)

        self.cnn = effnet.features

        self.channel_reducer = nn.Sequential(
            nn.Conv2d(1280, 512, kernel_size=1),
```

```

        nn.BatchNorm2d(512),
        nn.ReLU()
    )

    self.lstm = nn.LSTM(
        input_size=512,
        hidden_size=256,
        num_layers=2,
        bidirectional=True,
        batch_first=True
    )

    self.classifier = nn.Sequential(
        nn.Linear(512, 256),
        nn.ReLU(),
        nn.Dropout(0.5),
        nn.Linear(256, num_classes)
    )

    def forward(self, x):
        features = self.cnn(x)

        x = self.channel_reducer(features)
        bs, c, h, w = x.size()
        x = x.permute(0, 2, 3, 1).reshape(bs, h*w, c)

        lstm_out, (h_n, c_n) = self.lstm(x)
        last_hidden = torch.cat((h_n[-2], h_n[-1]), dim=1)

        return self.classifier(last_hidden)

model = EffNetLSTM(num_classes)
model.to('cuda')

# Loss and optimizer
import torch.optim as optim

wandb.watch(model, log="all")
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam([
    {'params': model.cnn.parameters(), 'lr': 1e-5},
    {'params': model.channel_reducer.parameters(), 'lr': 1e-4},
    {'params': model.lstm.parameters(), 'lr': 1e-4},
    {'params': model.classifier.parameters(), 'lr': 1e-4}
])

```

0.3 Training the model

```
[9]: # Train the model
num_epochs = 20

for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    train_bar = tqdm(train_loader, desc=f"Epoch {epoch+1}/{num_epochs}")
    for image_paths, labels in train_bar:
        image_tensors = torch.stack([get_image(image_path) for image_path in
↪image_paths])
        images = image_tensors.to('cuda')
        labels = labels.to('cuda')

        # Forward pass
        outputs = model(images)
        loss = criterion(outputs, labels)

        # Backward and optimize
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        train_bar.set_postfix(loss=loss.item())

    avg_train_loss = running_loss / len(train_loader)
    wandb.log({"epoch": epoch+1, "train_loss": avg_train_loss})

    # Validation Loop
    model.eval()
    val_loss = 0.0
    correct = 0
    total = 0
    with torch.no_grad():
        val_bar = tqdm(val_loader, desc="Validation")
        for image_paths, labels in val_bar:
            image_tensors = torch.stack([get_image(image_path) for image_path
↪in image_paths])
            image_tensors = image_tensors.to('cuda')
            labels = labels.to('cuda')
            outputs = model(image_tensors)
            loss = criterion(outputs, labels)
            val_loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
```

```

        correct += (predicted == labels.argmax(dim=1)).sum().item()
        val_bar.set_postfix(loss=loss.item())

    avg_val_loss = val_loss / len(val_loader)
    val_accuracy = 100 * correct / total
    wandb.log({"val_loss": avg_val_loss, "val_accuracy": val_accuracy})
    print(f"Epoch {epoch+1}/{num_epochs} - Train Loss: {avg_train_loss:.4f},  

    ↳ Val Loss: {avg_val_loss:.4f}, Val Accuracy: {val_accuracy:.2f}%")
    if (epoch%5==0):
        torch.save(model.state_dict(), f"effnet_rcnn_epoch_{epoch+1}_style.pth")
        torch.save(optimizer.state_dict(),  

    ↳ f"effnet_rcnn_optimizer_epoch_{epoch+1}_style.pth")

```

```

Epoch 1/20: 48%|          | 1718/3565 [10:12<11:26, 2.69it/s,
loss=2.23]Corrupt JPEG data: bad Huffman code
Epoch 1/20: 83%|          | 2958/3565 [17:25<03:57, 2.55it/s,
loss=2.17]Corrupt JPEG data: premature end of data segment
Epoch 1/20: 100%|         | 3565/3565 [20:57<00:00, 2.83it/s, loss=3.05]
Validation: 100%|         | 1527/1527 [06:48<00:00, 3.74it/s, loss=2.99]

Epoch 1/20 - Train Loss: 2.1877, Val Loss: 1.8370, Val Accuracy: 40.27%

Epoch 2/20: 5%|           | 180/3565 [01:02<17:15, 3.27it/s, loss=1.86]Corrupt
JPEG data: premature end of data segment
Epoch 2/20: 81%|          | 2905/3565 [16:46<04:00, 2.74it/s, loss=1.45]
Corrupt JPEG data: bad Huffman code
Epoch 2/20: 100%|         | 3565/3565 [20:40<00:00, 2.87it/s, loss=1.93]
Validation: 100%|         | 1527/1527 [06:40<00:00, 3.81it/s, loss=2.75]

Epoch 2/20 - Train Loss: 1.8250, Val Loss: 1.6779, Val Accuracy: 45.74%

Epoch 3/20: 95%|          | 3372/3565 [19:11<00:57, 3.34it/s, loss=1.84]
Corrupt JPEG data: bad Huffman code
Epoch 3/20: 98%|          | 3503/3565 [20:00<00:21, 2.92it/s, loss=1.39]
Corrupt JPEG data: premature end of data segment
Epoch 3/20: 100%|         | 3565/3565 [20:23<00:00, 2.91it/s, loss=2.12]
Validation: 100%|         | 1527/1527 [06:45<00:00, 3.76it/s, loss=2.6]

Epoch 3/20 - Train Loss: 1.6678, Val Loss: 1.5599, Val Accuracy: 48.49%

Epoch 4/20: 60%|          | 2137/3565 [12:19<06:57, 3.42it/s, loss=1.59]
Corrupt JPEG data: premature end of data segment
Epoch 4/20: 81%|          | 2894/3565 [17:01<03:16, 3.42it/s, loss=1.69]
Corrupt JPEG data: bad Huffman code
Epoch 4/20: 100%|         | 3565/3565 [21:09<00:00, 2.81it/s, loss=2.71]
Validation: 100%|         | 1527/1527 [06:55<00:00, 3.67it/s, loss=2.29]

Epoch 4/20 - Train Loss: 1.5504, Val Loss: 1.5222, Val Accuracy: 50.18%

Epoch 5/20: 43%|          | 1529/3565 [08:54<10:25, 3.25it/s, loss=1.84]
Corrupt JPEG data: bad Huffman code

```

Epoch 5/20: 60%| | 2141/3565 [12:28<09:31, 2.49it/s, loss=1.32]
Corrupt JPEG data: premature end of data segment
Epoch 5/20: 100%| | 3565/3565 [20:49<00:00, 2.85it/s, loss=1.22]
Validation: 100%| | 1527/1527 [06:25<00:00, 3.96it/s, loss=2.3]
Epoch 5/20 - Train Loss: 1.4688, Val Loss: 1.4633, Val Accuracy: 51.37%

Epoch 6/20: 31%| | 1093/3565 [06:07<14:14, 2.89it/s, loss=1.49]
Corrupt JPEG data: premature end of data segment
Epoch 6/20: 68%| | 2409/3565 [13:36<05:56, 3.24it/s, loss=2.02]
Corrupt JPEG data: bad Huffman code
Epoch 6/20: 100%| | 3565/3565 [20:15<00:00, 2.93it/s, loss=3.74]
Validation: 100%| | 1527/1527 [06:26<00:00, 3.95it/s, loss=2]
Epoch 6/20 - Train Loss: 1.3947, Val Loss: 1.4699, Val Accuracy: 51.94%

Epoch 7/20: 63%| | 2233/3565 [12:26<08:28, 2.62it/s, loss=1.05]
Corrupt JPEG data: bad Huffman code
Epoch 7/20: 90%| | 3194/3565 [18:23<02:38, 2.34it/s, loss=1.02]
Corrupt JPEG data: premature end of data segment
Epoch 7/20: 100%| | 3565/3565 [20:41<00:00, 2.87it/s, loss=6.95]
Validation: 100%| | 1527/1527 [06:55<00:00, 3.68it/s, loss=2.23]
Epoch 7/20 - Train Loss: 1.3154, Val Loss: 1.4130, Val Accuracy: 54.04%

Epoch 8/20: 62%| | 2206/3565 [12:48<06:47, 3.33it/s, loss=1.07]
Corrupt JPEG data: premature end of data segment
Epoch 8/20: 90%| | 3198/3565 [18:46<02:03, 2.97it/s, loss=1.39]
Corrupt JPEG data: bad Huffman code
Epoch 8/20: 100%| | 3565/3565 [20:56<00:00, 2.84it/s, loss=0.741]
Validation: 100%| | 1527/1527 [06:59<00:00, 3.64it/s, loss=1.74]
Epoch 8/20 - Train Loss: 1.2510, Val Loss: 1.3922, Val Accuracy: 54.58%

Epoch 9/20: 14%| | 491/3565 [03:02<17:49, 2.88it/s, loss=0.577]
Corrupt JPEG data: bad Huffman code
Epoch 9/20: 49%| | 1754/3565 [10:39<11:43, 2.57it/s, loss=1.27]
Corrupt JPEG data: premature end of data segment
Epoch 9/20: 100%| | 3565/3565 [21:50<00:00, 2.72it/s, loss=4.02]
Validation: 100%| | 1527/1527 [07:20<00:00, 3.47it/s, loss=1.58]
Epoch 9/20 - Train Loss: 1.1916, Val Loss: 1.3341, Val Accuracy: 56.07%

Epoch 10/20: 6%| | 224/3565 [01:19<21:12, 2.62it/s, loss=0.931]
Corrupt JPEG data: bad Huffman code
Epoch 10/20: 50%| | 1800/3565 [11:07<09:32, 3.08it/s, loss=0.867]
Corrupt JPEG data: premature end of data segment
Epoch 10/20: 100%| | 3565/3565 [21:35<00:00, 2.75it/s, loss=1.08]
Validation: 100%| | 1527/1527 [06:51<00:00, 3.71it/s, loss=2.02]
Epoch 10/20 - Train Loss: 1.1321, Val Loss: 1.3832, Val Accuracy: 55.46%

Epoch 11/20: 3%| | 96/3565 [00:34<19:56, 2.90it/s, loss=1.32]
Corrupt JPEG data: premature end of data segment

Epoch 11/20: 58%| | 2061/3565 [12:14<11:01, 2.27it/s, loss=1.68]
 Corrupt JPEG data: bad Huffman code
 Epoch 11/20: 100%| | 3565/3565 [21:02<00:00, 2.82it/s, loss=2.94]
 Validation: 100%| | 1527/1527 [06:34<00:00, 3.87it/s, loss=1.73]
 Epoch 11/20 - Train Loss: 1.0710, Val Loss: 1.3593, Val Accuracy: 56.45%

Epoch 12/20: 7%| | 254/3565 [01:24<15:14, 3.62it/s, loss=0.753]
 Corrupt JPEG data: bad Huffman code
 Epoch 12/20: 43%| | 1532/3565 [08:46<11:35, 2.92it/s, loss=1]
 Corrupt JPEG data: premature end of data segment
 Epoch 12/20: 100%| | 3565/3565 [20:30<00:00, 2.90it/s, loss=1.01]
 Validation: 100%| | 1527/1527 [06:36<00:00, 3.85it/s, loss=1.86]
 Epoch 12/20 - Train Loss: 1.0099, Val Loss: 1.3726, Val Accuracy: 56.44%

Epoch 13/20: 29%| | 1048/3565 [05:54<15:43, 2.67it/s, loss=0.536]
 Corrupt JPEG data: bad Huffman code
 Epoch 13/20: 86%| | 3059/3565 [17:26<03:39, 2.30it/s, loss=0.653]
 Corrupt JPEG data: premature end of data segment
 Epoch 13/20: 100%| | 3565/3565 [20:26<00:00, 2.91it/s, loss=2.25]
 Validation: 100%| | 1527/1527 [06:36<00:00, 3.85it/s, loss=1.95]
 Epoch 13/20 - Train Loss: 0.9635, Val Loss: 1.4224, Val Accuracy: 56.26%

Epoch 14/20: 67%| | 2381/3565 [13:40<06:36, 2.98it/s, loss=1.57]
 Corrupt JPEG data: bad Huffman code
 Epoch 14/20: 69%| | 2452/3565 [14:05<05:38, 3.29it/s, loss=1.16]
 Corrupt JPEG data: premature end of data segment
 Epoch 14/20: 100%| | 3565/3565 [20:53<00:00, 2.84it/s, loss=4.8]
 Validation: 100%| | 1527/1527 [07:13<00:00, 3.52it/s, loss=2.21]
 Epoch 14/20 - Train Loss: 0.9141, Val Loss: 1.3958, Val Accuracy: 56.37%

Epoch 15/20: 46%| | 1632/3565 [09:59<11:23, 2.83it/s, loss=0.893]
 Corrupt JPEG data: bad Huffman code
 Epoch 15/20: 97%| | 3462/3565 [21:02<00:39, 2.60it/s, loss=0.699]
 Corrupt JPEG data: premature end of data segment
 Epoch 15/20: 100%| | 3565/3565 [21:40<00:00, 2.74it/s, loss=8.51]
 Validation: 100%| | 1527/1527 [06:51<00:00, 3.71it/s, loss=2.95]
 Epoch 15/20 - Train Loss: 0.8691, Val Loss: 1.4523, Val Accuracy: 55.90%

Epoch 16/20: 26%| | 926/3565 [05:32<15:09, 2.90it/s, loss=0.872]
 Corrupt JPEG data: premature end of data segment
 Epoch 16/20: 65%| | 2309/3565 [13:51<07:32, 2.78it/s, loss=0.99]

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
Cell In[9], line 9
      7 train_bar = tqdm(train_loader, desc=f"Epoch {epoch+1}/{num_epochs}")
      8 for image_paths, labels in train_bar:
```

```

----> 9     image_tensors = torch.stack([get_image(image_path) for image_path i
↪image_paths])
      10     images = image_tensors.to('cuda')
      11     labels = labels.to('cuda')

```

Cell In[9], line 9, in <listcomp>(.0)

```

      7 train_bar = tqdm(train_loader, desc=f"Epoch {epoch+1}/{num_epochs}")
      8 for image_paths, labels in train_bar:
----> 9     image_tensors = torch.stack([get_image(image_path) for image_path i
↪image_paths])
      10     images = image_tensors.to('cuda')
      11     labels = labels.to('cuda')

```

Cell In[7], line 5, in get_image(image_path, image_size)

```

      3 def get_image(image_path, image_size=224):
      4     try:
----> 5         img = cv2.imread('./wikiart/' + image_path)
      6         if img is None:
      7             raise ValueError(f"Image not loaded: ./wikiart/{image_path} ")

```

KeyboardInterrupt: