

similarity

March 25, 2025

0.1 Downloading the national gallery of art open data from kaggle

```
[ ]: import kagglehub

path = kagglehub.dataset_download("peacehegemony/
↳the-national-gallery-of-art-open-data-program")

print("Path to dataset files:", path)
```

Path to dataset files: /home/ameya/.cache/kagglehub/datasets/peacehegemony/the-national-gallery-of-art-open-data-program/versions/1

```
[2]: path = path + "/opendata-main/data/"
```

0.2 Joining table based on appropriate columns to get data

```
[ ]: import pandas as pd

objects_df = pd.read_csv(path + "objects.csv")
mediarel_df = pd.read_csv(path + "media_relationships.csv")
media_items_df = pd.read_csv(path + "media_items.csv")

obj_mediarel_df = objects_df.merge(
    mediarel_df,
    left_on="objectid",
    right_on="relatedid",
    how="inner"
)

merged_df = obj_mediarel_df.merge(
    media_items_df,
    on="mediaid",
    how="inner"
)

print("Merged DataFrame shape:", merged_df.shape)
print(merged_df.head())
```

/tmp/ipykernel_108246/2913492502.py:3: DtypeWarning: Columns (29) have mixed types. Specify dtype option on import or set low_memory=False.

```
objects_df = pd.read_csv(path + "objects.csv")
```

Merged DataFrame shape: (2886, 47)

	objectid	accessioned	accessionnum	locationid	title_x
0	61	1	1937.1.54	NaN	The Lacemaker \
1	62	1	1937.1.55	NaN	The Smiling Girl
2	62	1	1937.1.55	NaN	The Smiling Girl
3	2718	1	1942.9.1839	NaN	Oeuvres poissardes
4	62	1	1937.1.55	NaN	The Smiling Girl

	displaydate	beginyear	endyear	visualbrowsertimespan
0	c. 1925	1925.0	1925.0	1901 to 1925 \
1	c. 1925	1925.0	1925.0	1901 to 1925
2	c. 1925	1925.0	1925.0	1901 to 1925
3	published 1796	1796.0	1796.0	1776 to 1800
4	c. 1925	1925.0	1925.0	1901 to 1925

	medium	language
0	oil on canvas	en \
1	oil on canvas	en
2	oil on canvas	en
3	1 vol: ill: 4 color stipple engravings	en
4	oil on canvas	en

	thumbnailurl
0	https://www.nga.gov/content/dam/ngaweb/audio-v... \
1	https://www.nga.gov/content/dam/ngaweb/audio-v...
2	https://www.nga.gov/content/dam/ngaweb/audio-v...
3	https://www.nga.gov/content/dam/ngaweb/audio-v...
4	https://www.nga.gov/content/dam/ngaweb/audio-v...

	playurl
0	https://w.soundcloud.com/player/?url=https%3A%... \
1	https://w.soundcloud.com/player/?url=https%3A%...
2	https://w.soundcloud.com/player/?url=https%3A%...
3	https://w.soundcloud.com/player/?url=https%3A%...
4	https://players.brightcove.net/1191289016001/d...

	downloadurl
0	https://api.soundcloud.com/tracks/78345258/dow... \
1	https://api.soundcloud.com/tracks/78345258/dow...
2	https://api.soundcloud.com/tracks/475697055/do...
3	https://api.soundcloud.com/tracks/475697055/do...
4	NaN

keywords

```

0     vermeer, han van meegeren, forgery \
1     vermeer, han van meegeren, forgery
2 jecmen, rosenwald, prints, drawings,
3 jecmen, rosenwald, prints, drawings,
4 jecmen, rosenwald, prints, drawings,

                                tags
0 ngaweb:audio-video/audio,ngaweb:audio-video/au... \
1 ngaweb:audio-video/audio,ngaweb:audio-video/au...
2 ngaweb:constituents/6/2/Constituent_62,ngaweb:...
3 ngaweb:constituents/6/2/Constituent_62,ngaweb:...
4 ngaweb:audio-video/podcast-video,ngaweb:audio-...

                                imageurl                presentationdate
0 https://www.nga.gov/content/dam/ngaweb/audio-v... 2009-01-11 00:00:00-05 \
1 https://www.nga.gov/content/dam/ngaweb/audio-v... 2009-01-11 00:00:00-05
2 https://www.nga.gov/content/dam/ngaweb/audio-v... 2018-03-16 00:00:00-04
3 https://www.nga.gov/content/dam/ngaweb/audio-v... 2018-03-16 00:00:00-04
4 https://www.nga.gov/content/dam/ngaweb/audio-v... 2018-03-16 00:00:00-04

                                releasedate                lastmodified
0 2009-01-20 00:00:00-05 2014-10-10 12:22:21-04
1 2009-01-20 00:00:00-05 2014-10-10 12:22:21-04
2 2018-07-24 00:00:00-04 2019-04-09 15:14:14-04
3 2018-07-24 00:00:00-04 2019-04-09 15:14:14-04
4 2018-07-24 00:00:00-04 2019-04-09 15:14:30-04

```

[5 rows x 47 columns]

0.3 Creating a dictionary with objectid as the key and related imageurls as the values

```

[4]: images_dict = {}

for idx, row in merged_df.iterrows():
    obj_id = row["objectid"]
    img_path = row["imageurl"]

    if obj_id not in images_dict:
        images_dict[obj_id] = []

    images_dict[obj_id].append(img_path)

# test_ids = list(images_dict.keys())
# for test_id in test_ids:
#     print(f"Object ID: {test_id}")
#     print("Image paths:", images_dict[test_id])

```

```
print(len(images_dict))
```

807

```
[5]: import os
import random
from PIL import Image
import torch
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as transforms
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import requests
from io import BytesIO
```

0.4 Loading image from the given url

```
[ ]: issues = 0
total = 0
def load_image_from_url(url):
    global issues, total
    try:
        total += 1
        response = requests.get(url, timeout=10)
        response.raise_for_status()
        img = Image.open(BytesIO(response.content)).convert("RGB")
        return img
    except Exception as e:
        issues += 1
        return Image.new("RGB", (224, 224), (0, 0, 0))
```

0.5 Creating a dataset for siamese networks (consisting of negative and positive samples)

```
[ ]: class SiameseDataset(Dataset):
    def __init__(self, images_dict, transform=None, target_size=(224, 224)):
        self.images_dict = images_dict
        self.object_ids = list(images_dict.keys())
        self.transform = transform
        self.target_size = target_size
        self.pairs = []
        self.labels = []
        self._create_pairs()

    def _create_pairs(self):
```

```

        for obj_id in self.object_ids:
            urls = self.images_dict[obj_id]
            if len(urls) < 2:
                continue
            for i in range(len(urls)):
                for j in range(i+1, len(urls)):
                    self.pairs.append((urls[i], urls[j]))
                    self.labels.append(1)

        num_positive = len(self.labels)
        neg_pairs = 0
        while neg_pairs < num_positive:
            id1, id2 = random.sample(self.object_ids, 2)
            if len(self.images_dict[id1]) == 0 or len(self.images_dict[id2]) == 0:
                continue
            url1 = random.choice(self.images_dict[id1])
            url2 = random.choice(self.images_dict[id2])
            self.pairs.append((url1, url2))
            self.labels.append(0)
            neg_pairs += 1

    def __len__(self):
        return len(self.pairs)

    def __getitem__(self, idx):
        url1, url2 = self.pairs[idx]
        img1 = load_image_from_url(url1)
        img2 = load_image_from_url(url2)
        if self.transform:
            img1 = self.transform(img1)
            img2 = self.transform(img2)
        label = torch.tensor(self.labels[idx], dtype=torch.float32)
        return img1, img2, label

# Define transformations for the images
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                          std=[0.229, 0.224, 0.225])
])

```

```

[8]: dataset = SiameseDataset(images_dict, transform=transform)
     # split into test and train
     train_size = int(0.8 * len(dataset))
     test_size = len(dataset) - train_size

```

```

train_dataset, test_dataset = torch.utils.data.random_split(dataset,
    ↪[train_size, test_size])

train_loader = DataLoader(train_dataset, batch_size=32,
    ↪shuffle=True,num_workers=4)
test_loader = DataLoader(test_dataset, batch_size=32,
    ↪shuffle=False,num_workers=4)

```

```

[ ]: class SiameseNetwork(nn.Module):
    def __init__(self):
        super(SiameseNetwork, self).__init__()
        self.cnn = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2),
            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2),
            nn.Conv2d(64, 128, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2)
        )
        self.fc = nn.Sequential(
            nn.Linear(128 * 28 * 28, 512),
            nn.ReLU(inplace=True),
            nn.Linear(512, 128)
        )

    def forward_once(self, x):
        output = self.cnn(x)
        output = output.view(output.size(0), -1)
        output = self.fc(output)
        return output

    def forward(self, input1, input2):
        output1 = self.forward_once(input1)
        output2 = self.forward_once(input2)
        return output1, output2

```

0.6 Contrastive loss function for training

```

[ ]: class ContrastiveLoss(nn.Module):
    def __init__(self, margin=1.0):
        super(ContrastiveLoss, self).__init__()
        self.margin = margin

```

```

def forward(self, output1, output2, label):
    euclidean_distance = F.pairwise_distance(output1, output2)
    loss_contrastive = torch.mean(
        label * torch.pow(euclidean_distance, 2) +
        (1 - label) * torch.pow(torch.clamp(self.margin -
↪euclidean_distance, min=0.0), 2)
    )
    return loss_contrastive

```

0.7 Model declaration

```

[11]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = SiameseNetwork().to(device)
criterion = ContrastiveLoss(margin=1.0)
optimizer = optim.Adam(model.parameters(), lr=1e-3)

```

```

[12]: from tqdm import tqdm

```

0.8 Training the model

```

[13]: num_epochs = 11

for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    train_bar = tqdm(train_loader, desc=f"Epoch {epoch+1}/{num_epochs} -
↪Training", leave=False)
    for img1, img2, label in train_bar:
        img1, img2, label = img1.to(device), img2.to(device), label.to(device)

        optimizer.zero_grad()
        output1, output2 = model(img1, img2)
        loss = criterion(output1, output2, label)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        train_bar.set_postfix(loss=f"{loss.item():.4f}")

    avg_loss = running_loss / len(train_loader)
    print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {avg_loss:.4f}")

    model.eval()
    test_loss = 0.0
    test_bar = tqdm(test_loader, desc=f"Epoch {epoch+1}/{num_epochs} -
↪Testing", leave=False)

```

```

with torch.no_grad():
    for img1, img2, label in test_bar:
        img1, img2, label = img1.to(device), img2.to(device), label.
        ↪to(device)
        output1, output2 = model(img1, img2)
        loss = criterion(output1, output2, label)
        test_loss += loss.item()
        test_bar.set_postfix(loss=f"{loss.item():.4f}")
avg_test_loss = test_loss / len(test_loader)
print(f"Test Loss: {avg_test_loss:.4f}")
if(epoch % 5 == 0):
    torch.save(model.state_dict(), f"siamese_model_epoch{epoch}.pt")

print("Training complete.")

```

Epoch 1/11 - Training: 0% | 0/742 [00:00<?, ?it/s]

Epoch [1/11], Loss: 0.3343

Test Loss: 0.0724

Epoch [2/11], Loss: 0.0443

Test Loss: 0.0493

Epoch [3/11], Loss: 0.0306

Test Loss: 0.0361

Epoch [4/11], Loss: 0.0248

Test Loss: 0.0300

Epoch [5/11], Loss: 0.0222

Test Loss: 0.0277

Epoch [6/11], Loss: 0.0210

Test Loss: 0.0253

Epoch [7/11], Loss: 0.0188

Test Loss: 0.0286

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
Cell In[13], line 8
      6 # Wrap the training dataloader with tqdm directly
      7 train_bar = tqdm(train_loader, desc=f"Epoch {epoch+1}/{num_epochs} - Training", leave=False)
---->  8 for img1, img2, label in train_bar:
      9     img1, img2, label = img1.to(device), img2.to(device), label.
      9     to(device)
     11     optimizer.zero_grad()

File ~/.local/lib/python3.10/site-packages/tqdm/std.py:1178, in tqdm.
     1175 time = self._time
     1177 try:
-> 1178     for obj in iterable:
     1179         yield obj
     1180         # Update and possibly print the progressbar.
     1181         # Note: does not call self.update(1) for speed optimisation.

File ~/.local/lib/python3.10/site-packages/torch/utils/data/dataloader.py:708, in _BaseDataLoaderIter.__next__(self)
     705 if self._sampler_iter is None:
     706     # TODO(https://github.com/pytorch/pytorch/issues/76750)
     707     self._reset() # type: ignore[call-arg]
--> 708 data = self._next_data()
     709 self._num_yielded += 1
     710 if (
     711     self._dataset_kind == _DatasetKind.Iterable
     712     and self._IterableDataset_len_called is not None
     713     and self._num_yielded > self._IterableDataset_len_called
     714 ):

File ~/.local/lib/python3.10/site-packages/torch/utils/data/dataloader.py:1458, in _MultiProcessingDataLoaderIter._next_data(self)
    1455     return self._process_data(data)
```

```

1457 assert not self._shutdown and self._tasks_outstanding > 0
-> 1458 idx, data = self._get_data()
1459 self._tasks_outstanding -= 1
1460 if self._dataset_kind == _DatasetKind.Iterable:
1461     # Check for _IterableDatasetStopIteration

File ~/.local/lib/python3.10/site-packages/torch/utils/data/dataloader.py:1420,
in _MultiProcessingDataLoaderIter._get_data(self)
    1416     # In this case, `self._data_queue` is a `queue.Queue`,. But we don't
    1417     # need to call `.task_done()` because we don't use `.join()`.
    1418 else:
    1419     while True:
-> 1420         success, data = self._try_get_data()
    1421         if success:
    1422             return data

File ~/.local/lib/python3.10/site-packages/torch/utils/data/dataloader.py:1251,
in _MultiProcessingDataLoaderIter._try_get_data(self, timeout)
    1238 def _try_get_data(self, timeout=_utils.MP_STATUS_CHECK_INTERVAL):
    1239     # Tries to fetch data from `self._data_queue` once for a given
-> timeout.
    1240     # This can also be used as inner loop of fetching without timeout,
-> with
    1241     (...)
    1242     # Returns a 2-tuple:
    1243     # (bool: whether successfully get data, any: data if successful,
-> else None)
    1244     try:
-> 1251         data = self._data_queue.get(timeout=timeout)
    1252         return (True, data)
    1253     except Exception as e:
    1254         # At timeout and error, we manually check whether any worker ha
    1255         # failed. Note that this is the only mechanism for Windows to
-> detect
    1256         # worker failures.

File /usr/lib/python3.10/multiprocessing/queues.py:113, in Queue.get(self,
-> block, timeout)
    111 if block:
    112     timeout = deadline - time.monotonic()
--> 113     if not self._poll(timeout):
    114         raise Empty
    115 elif not self._poll():

File /usr/lib/python3.10/multiprocessing/connection.py:257, in _ConnectionBase.
-> poll(self, timeout)
    255 self._check_closed()
    256 self._check_readable()

```

```
--> 257 return self._poll(timeout)
```

File /usr/lib/python3.10/multiprocessing/connection.py:424, in Connection.

```
↪ _poll(self, timeout)
    423 def _poll(self, timeout):
--> 424     r = wait([self], timeout)
    425     return bool(r)
```

File /usr/lib/python3.10/multiprocessing/connection.py:931, in wait(object_list, timeout)

```
↪ timeout)
    928     deadline = time.monotonic() + timeout
    930 while True:
--> 931     ready = selector.select(timeout)
    932     if ready:
    933         return [key.fileobj for (key, events) in ready]
```

File /usr/lib/python3.10/selectors.py:416, in _PollLikeSelector.select(self, timeout)

```
↪ timeout)
    414 ready = []
    415 try:
--> 416     fd_event_list = self._selector.poll(timeout)
    417 except InterruptedError:
    418     return ready
```

KeyboardInterrupt:

0.9 Dataset creation (consisting of anchor, positive and negative images)

```
[14]: class TripletDataset(Dataset):
    def __init__(self, images_dict, transform=None, target_size=(224, 224)):
        self.images_dict = images_dict
        self.object_ids = list(images_dict.keys())
        self.transform = transform
        self.target_size = target_size
        self.all_images = [(obj_id, url) for obj_id, urls in images_dict.
↪ items() for url in urls]

    def __len__(self):
        return len(self.all_images)

    def __getitem__(self, idx):
        anchor_obj_id, anchor_url = self.all_images[idx]
        anchor_img = load_image_from_url(anchor_url)
        if self.transform:
            anchor_img = self.transform(anchor_img)
```

```

        positive_urls = self.images_dict[anchor_obj_id]
        if len(positive_urls) < 2:
            positive_img = anchor_img.clone()
        else:
            candidate_urls = [url for url in positive_urls if url != anchor_url]
            if candidate_urls:
                pos_url = random.choice(candidate_urls)
            else:
                pos_url = anchor_url
            positive_img = load_image_from_url(pos_url)
            if self.transform:
                positive_img = self.transform(positive_img)

        negative_obj_id = random.choice([oid for oid in self.object_ids if oid !=
↪ anchor_obj_id])
        neg_url = random.choice(self.images_dict[negative_obj_id])
        negative_img = load_image_from_url(neg_url)
        if self.transform:
            negative_img = self.transform(negative_img)

        return anchor_img, positive_img, negative_img

# Define image transformations
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                           std=[0.229, 0.224, 0.225])
])

```

```

[15]: triplet_dataset = TripletDataset(images_dict, transform=transform)

```

```

# Split into train and test sets
train_size = int(0.8 * len(triplet_dataset))
test_size = len(triplet_dataset) - train_size
triplet_train_dataset, triplet_test_dataset = torch.utils.data.
↪ random_split(triplet_dataset, [train_size, test_size])
triplet_train_loader = DataLoader(triplet_train_dataset, batch_size=16,
↪ shuffle=True, num_workers=4)
triplet_test_loader = DataLoader(triplet_test_dataset, batch_size=16,
↪ shuffle=False, num_workers=4)

```

```

[16]: class FeatureExtractor(nn.Module):
        def __init__(self):
            super(FeatureExtractor, self).__init__()
            self.cnn = nn.Sequential(
                nn.Conv2d(3, 32, kernel_size=3, padding=1),

```

```

        nn.ReLU(inplace=True),
        nn.MaxPool2d(2),
        nn.Conv2d(32, 64, kernel_size=3, padding=1),
        nn.ReLU(inplace=True),
        nn.MaxPool2d(2),
        nn.Conv2d(64, 128, kernel_size=3, padding=1),
        nn.ReLU(inplace=True),
        nn.MaxPool2d(2)
    )
    self.fc = nn.Sequential(
        nn.Linear(128 * 28 * 28, 512),
        nn.ReLU(inplace=True),
        nn.Linear(512, 128)
    )

    def forward(self, x):
        x = self.cnn(x)
        x = x.view(x.size(0), -1)
        x = self.fc(x)
        return x

```

```

[17]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = FeatureExtractor().to(device)
margin = 1.0
criterion = nn.TripletMarginLoss(margin=margin, p=2)
optimizer = optim.Adam(model.parameters(), lr=1e-3)

```

0.10 Training model with triplet loss function

```

[18]: num_epochs = 11
for epoch in range(num_epochs):
    running_loss = 0.0
    for batch_idx, (anchor, positive, negative) in in
        ↪ enumerate(tqdm(triplet_train_loader, desc=f"Epoch {epoch+1}/{num_epochs}")):
            anchor, positive, negative = anchor.to(device), positive.to(device), ↪
            ↪ negative.to(device)
            optimizer.zero_grad()

            # Compute embeddings
            anchor_out = model(anchor)
            positive_out = model(positive)
            negative_out = model(negative)

            loss = criterion(anchor_out, positive_out, negative_out)
            loss.backward()
            optimizer.step()

```

```

        running_loss += loss.item()
    avg_loss = running_loss / len(triplet_train_loader)
    print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {avg_loss:.4f}")
    if(epoch % 5 == 0):
        torch.save(model.state_dict(), f"triplet_model_epoch{epoch}.pt")

    # Evaluate the model on the test set
    model.eval()
    test_loss = 0.0
    with torch.no_grad():
        for anchor, positive, negative in tqdm(triplet_test_loader,
        desc=f"Epoch {epoch+1}/{num_epochs} - Testing"):
            anchor, positive, negative = anchor.to(device), positive.
            to(device), negative.to(device)
            anchor_out = model(anchor)
            positive_out = model(positive)
            negative_out = model(negative)
            loss = criterion(anchor_out, positive_out, negative_out)
            test_loss += loss.item()
    avg_test_loss = test_loss / len(triplet_test_loader)
    print(f"Test Loss: {avg_test_loss:.4f}")

print("Training complete.")

```

```

Epoch 1/11: 100%|          | 145/145 [08:12<00:00,  3.40s/it]
Epoch [1/11], Loss: 0.5880
Epoch 1/11 - Testing: 100%|          | 37/37 [01:43<00:00,  2.81s/it]
Test Loss: 0.5996
Epoch 2/11: 100%|          | 145/145 [05:28<00:00,  2.26s/it]
Epoch [2/11], Loss: 0.5261
Epoch 2/11 - Testing: 100%|          | 37/37 [01:19<00:00,  2.14s/it]
Test Loss: 0.7084
Epoch 3/11: 100%|          | 145/145 [04:37<00:00,  1.91s/it]
Epoch [3/11], Loss: 0.5170
Epoch 3/11 - Testing: 100%|          | 37/37 [01:08<00:00,  1.85s/it]
Test Loss: 0.4872
Epoch 4/11: 100%|          | 145/145 [04:24<00:00,  1.82s/it]
Epoch [4/11], Loss: 0.5462

```

Epoch 4/11 - Testing: 100%| | 37/37 [01:07<00:00, 1.82s/it]
Test Loss: 0.6398
Epoch 5/11: 100%| | 145/145 [04:29<00:00, 1.86s/it]
Epoch [5/11], Loss: 0.5368
Epoch 5/11 - Testing: 100%| | 37/37 [01:10<00:00, 1.91s/it]
Test Loss: 0.4366
Epoch 6/11: 100%| | 145/145 [04:31<00:00, 1.87s/it]
Epoch [6/11], Loss: 0.5100
Epoch 6/11 - Testing: 100%| | 37/37 [01:07<00:00, 1.82s/it]
Test Loss: 0.5094
Epoch 7/11: 100%| | 145/145 [04:02<00:00, 1.67s/it]
Epoch [7/11], Loss: 0.4166
Epoch 7/11 - Testing: 100%| | 37/37 [01:03<00:00, 1.72s/it]
Test Loss: 0.3756
Epoch 8/11: 100%| | 145/145 [04:15<00:00, 1.76s/it]
Epoch [8/11], Loss: 0.3526
Epoch 8/11 - Testing: 100%| | 37/37 [01:01<00:00, 1.66s/it]
Test Loss: 0.3378
Epoch 9/11: 100%| | 145/145 [04:03<00:00, 1.68s/it]
Epoch [9/11], Loss: 0.3199
Epoch 9/11 - Testing: 100%| | 37/37 [01:04<00:00, 1.75s/it]
Test Loss: 0.3420
Epoch 10/11: 100%| | 145/145 [04:07<00:00, 1.71s/it]
Epoch [10/11], Loss: 0.2642
Epoch 10/11 - Testing: 100%| | 37/37 [01:02<00:00, 1.69s/it]
Test Loss: 0.2602
Epoch 11/11: 100%| | 145/145 [04:19<00:00, 1.79s/it]
Epoch [11/11], Loss: 0.2014
Epoch 11/11 - Testing: 100%| | 37/37 [01:12<00:00, 1.96s/it]
Test Loss: 0.2251
Training complete.

0.11 Evaluating the models

```
[20]: model1 = FeatureExtractor().to(device)
model1.load_state_dict(torch.load("triplet_model_epoch10.pt"))

model2 = SiameseNetwork().to(device)
model2.load_state_dict(torch.load("siamese_model_epoch5.pt"))
```

```
[20]: <All keys matched successfully>
```

```
[21]: ## Using the validation set to test the models

def evaluate_model(model, dataloader):
    model.eval()
    embeddings1 = []
    embeddings2 = []
    labels = []
    with torch.no_grad():
        for img1, img2, label in tqdm(dataloader):
            img1, img2 = img1.to(device), img2.to(device)
            output1, output2 = model(img1, img2)
            embeddings1.append(output1)
            embeddings2.append(output2)
            labels.append(label)
    embeddings1 = torch.cat(embeddings1)
    embeddings2 = torch.cat(embeddings2)
    labels = torch.cat(labels)
    return embeddings1, embeddings2, labels

siamese_embeddings1, siamese_embeddings2, siamese_labels = \
    ↪ evaluate_model(model2, test_loader)

## Calculate the average cosine similarity for positive and negative pairs ↪
↪ respectively

def calculate_cosine_similarity(embeddings1, embeddings2):
    cosine_similarity = nn.CosineSimilarity(dim=1)
    return cosine_similarity(embeddings1, embeddings2)

positive_indices = siamese_labels == 1
negative_indices = siamese_labels == 0

positive_similarities = \
    ↪ calculate_cosine_similarity(siamese_embeddings1[positive_indices], \
    ↪ siamese_embeddings2[positive_indices])
negative_similarities = \
    ↪ calculate_cosine_similarity(siamese_embeddings1[negative_indices], \
    ↪ siamese_embeddings2[negative_indices])
```



```

print(f"Average Cosine Similarity for Positive Pairs: {positive_similarities.
↪mean().item():.4f}")
print(f"Average Cosine Similarity for Negative Pairs: {negative_similarities.
↪mean().item():.4f}")

```

100%| | 186/186 [07:09<00:00, 2.31s/it]

Average Cosine Similarity for Positive Pairs: 0.9815

Average Cosine Similarity for Negative Pairs: 0.2679

```

[23]: cosine_similarities_similar = []
      cosine_similarities_dissimilar = []

      with torch.no_grad():
          for anchor, positive, negative in tqdm(triplet_test_loader):
              anchor, positive, negative = anchor.to(device), positive.to(device), ↪
              ↪negative.to(device)
              anchor_out = model1(anchor)
              positive_out = model1(positive)
              negative_out = model1(negative)

              cos_sim_similar = F.cosine_similarity(anchor_out, positive_out, dim=1)
              cos_sim_dissimilar = F.cosine_similarity(anchor_out, negative_out, ↪
              ↪dim=1)

              cosine_similarities_similar.append(cos_sim_similar)
              cosine_similarities_dissimilar.append(cos_sim_dissimilar)

      avg_cos_sim_similar = torch.cat(cosine_similarities_similar).mean().item()
      avg_cos_sim_dissimilar = torch.cat(cosine_similarities_dissimilar).mean().item()

      print(f"Average Cosine Similarity (Anchor-Positive, Similar): ↪
      ↪{avg_cos_sim_similar:.4f}")
      print(f"Average Cosine Similarity (Anchor-Negative, Dissimilar): ↪
      ↪{avg_cos_sim_dissimilar:.4f}")

```

100%| | 37/37 [00:58<00:00, 1.58s/it]

Average Cosine Similarity (Anchor-Positive, Similar): 0.7788

Average Cosine Similarity (Anchor-Negative, Dissimilar): 0.2432